



# Introdução à Programação Paralela em Clusters Heterogêneos

Silvio Stanzani, Jefferson Fialho , Raphael Cóbe , Rogério Iope  
Núcleo de Computação Científica – UNESP

[silvio, fialho, rmcobe, rogerio]@ncc.unesp.br

**ERAD-SP 2018**

# UNESP Center for Scientific Computing

- Consolidates scientific computing resources for São Paulo State University (UNESP) researchers
  - It mainly uses Grid computing paradigm
- Main users
  - UNESP researchers, students, and software developers
  - SPRACE (São Paulo Research and Analysis Center)
  - physicists and students
- Caltech, Fermilab, CERN
- São Paulo CMS Tier-2 Facility

# Agenda

- Parallel Architectures
- Parallelism Levels
- Xeon and Xeon Phi
- Heterogeneous Cluster @ NCC
- Offload Programming
- Optimization Examples
- Hands-on

# Agenda

- **Parallel Architectures**
- Parallelism Levels
- Xeon and Xeon Phi
- Heterogeneous Cluster @ NCC
- Offload Programming
- Optimization Examples
- Hands-on


# Parallel Processing

- ❑ A parallel computer is a computer system that uses multiple processing elements simultaneously in a cooperative manner to solve a computational problem
- ❑ Parallel processing includes techniques and technologies that make it possible to compute in parallel
  - ❑ Hardware, networks, operating systems, parallel libraries, languages, compilers, algorithms, tools, ...
- ❑ Parallel computing is an evolution of serial computing
  - ❑ Parallelism is natural
  - ❑ Computing problems differ in level / type of parallelism

# Parallelism

- Concurrency = Different units of work that can be executed in parallel
  - Ex: loop iterations
- Parallelism = concurrency + parallel hardware
  - Find concurrent execution opportunities
  - Develop application to execute in parallel
  - Run application on parallel hardware
- Motivations for parallelism
  - Faster time to solution
  - Solve bigger computing problems
  - Effective use of machine resources
- Serial machines have inherent limitations
  - Processor speed, memory bottlenecks, ...

# Homogeneous Parallel Architectures

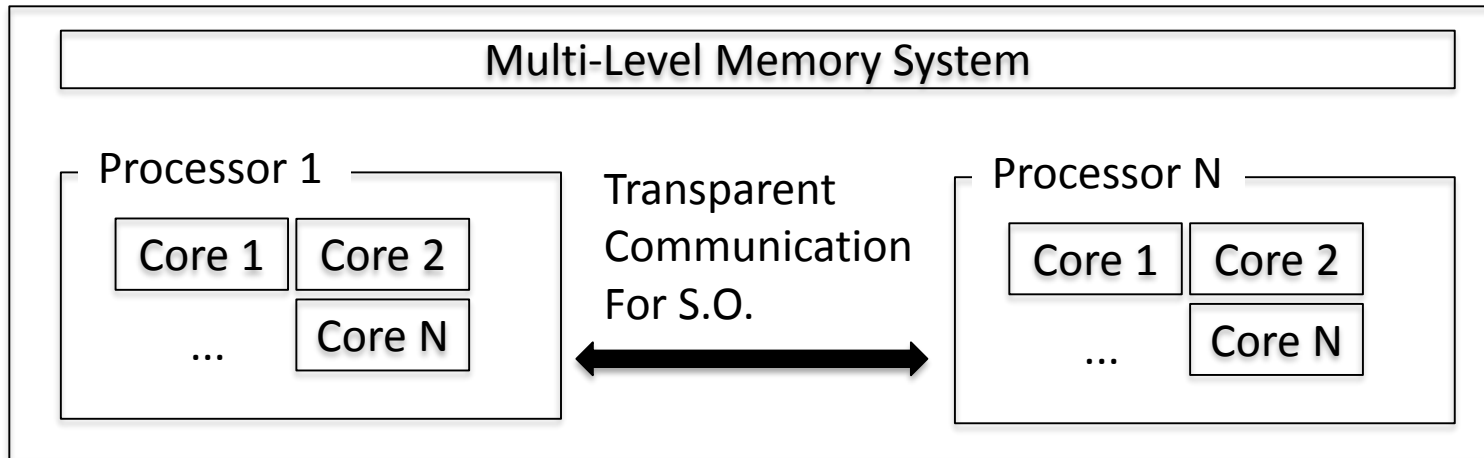


- **First initiatives:**

- Identical Nodes interconnected;
- Parallel architectures is deployed as a group of serial computers;


# Heterogeneous Parallel Architectures

## Multi-Core Node



## Multi-Core Architecture


- Multi Processors
- Multiple Cores
- **NUMA** (Non-Uniform Memory Access)





# Heterogeneous Parallel Architectures

- Several families and architectures:
- Hybrid Parallel nodes
  - Coprocessors and accelerators;
- Multi-level parallelism:
  - Processing core;
  - Chip multiprocessor;
  - Computing node;
  - Computing cluster;



# Agenda

- Parallel Architectures
- **Parallelism Levels**
- Xeon and Xeon Phi
- Heterogeneous Cluster @ NCC
- Offload Programming
- Optimization Examples
- Hands-on

# Exploiting the parallel universe

## Instruction Level Parallelism

- Single thread (ST) performance
- Automatically exposed by HW/tools
- Effectively limited to a few instructions

## Data Level Parallelism

- Single thread (ST) performance
- Exposed by tools and programming models
- Operate on 4/8/16 elements at a time

**Vectorization**

## Task Level Parallelism

- Multi thread/task (MT) performance
- Exposed by programming models
- Execute tens/hundreds/thousands task concurrently

**OpenMP**

## Process Level Parallelism

- Multi Process (MP) performance
- Exposed by programming models
- Execute tens/hundreds/thousands of process concurrently across several nodes

**MPI**




# Vectorization


- Vectorization
  - Loading data into cache accordingly;
  - Store elements on SIMD registers or vectors;
  - Apply the same operation to a set of Data at the same time;
  - Iterations need to be independent;
  - Usually on inner loops.

## Scalar loop

```
for (int i = 0; i < N; i++)  
    c[i] = a[i] + b[i];
```




# OpenMP



# Offload

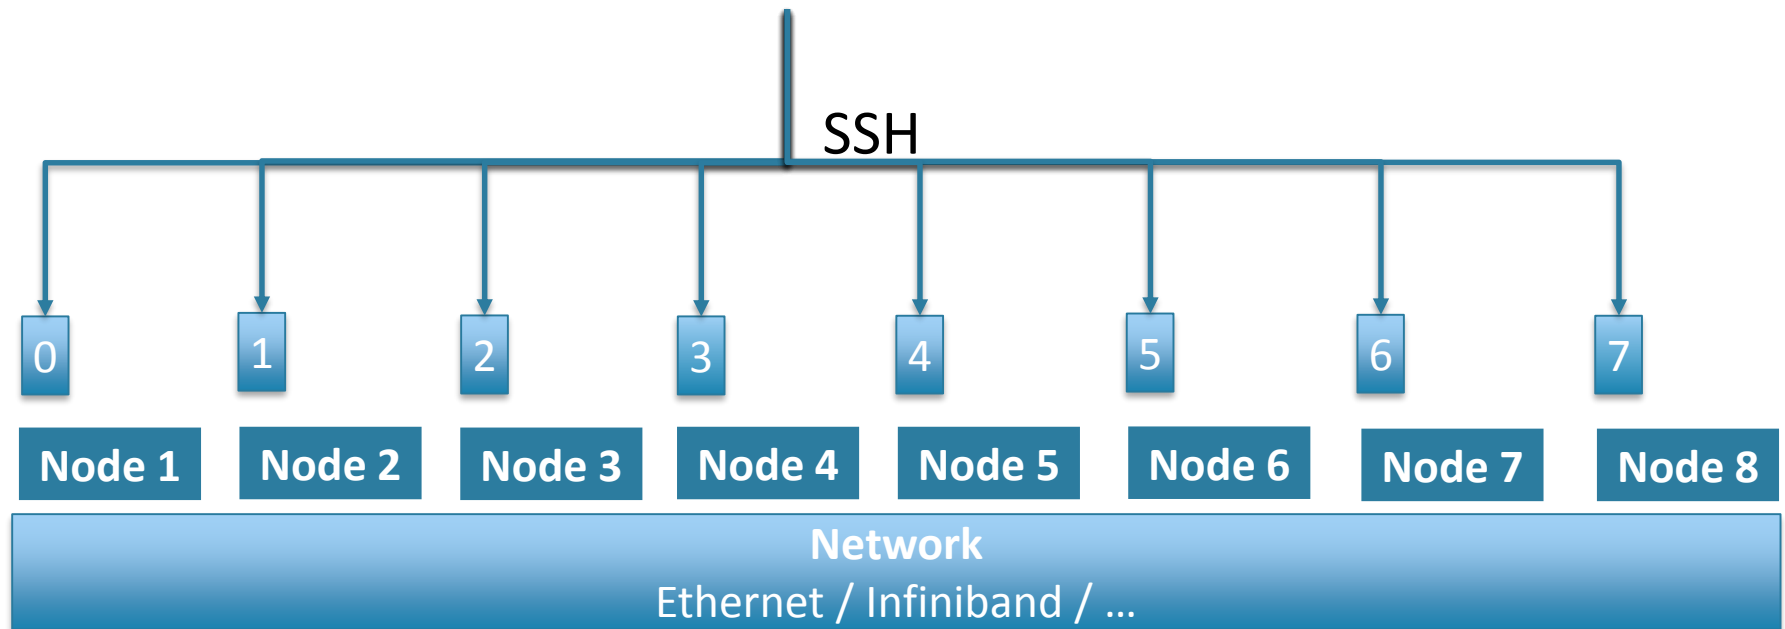
- Regions of code is marked to be executed in a device attached to the main processor
  - Communication using PCI Express or Network
- Main Processor and Device has independent Memory system
  - Developer need to control data transfer of data structures needed by application
  - Can Present Delay



# Message Passing Interface (MPI)

Execution of a MPI application across several nodes using SSH (Secure Shell)

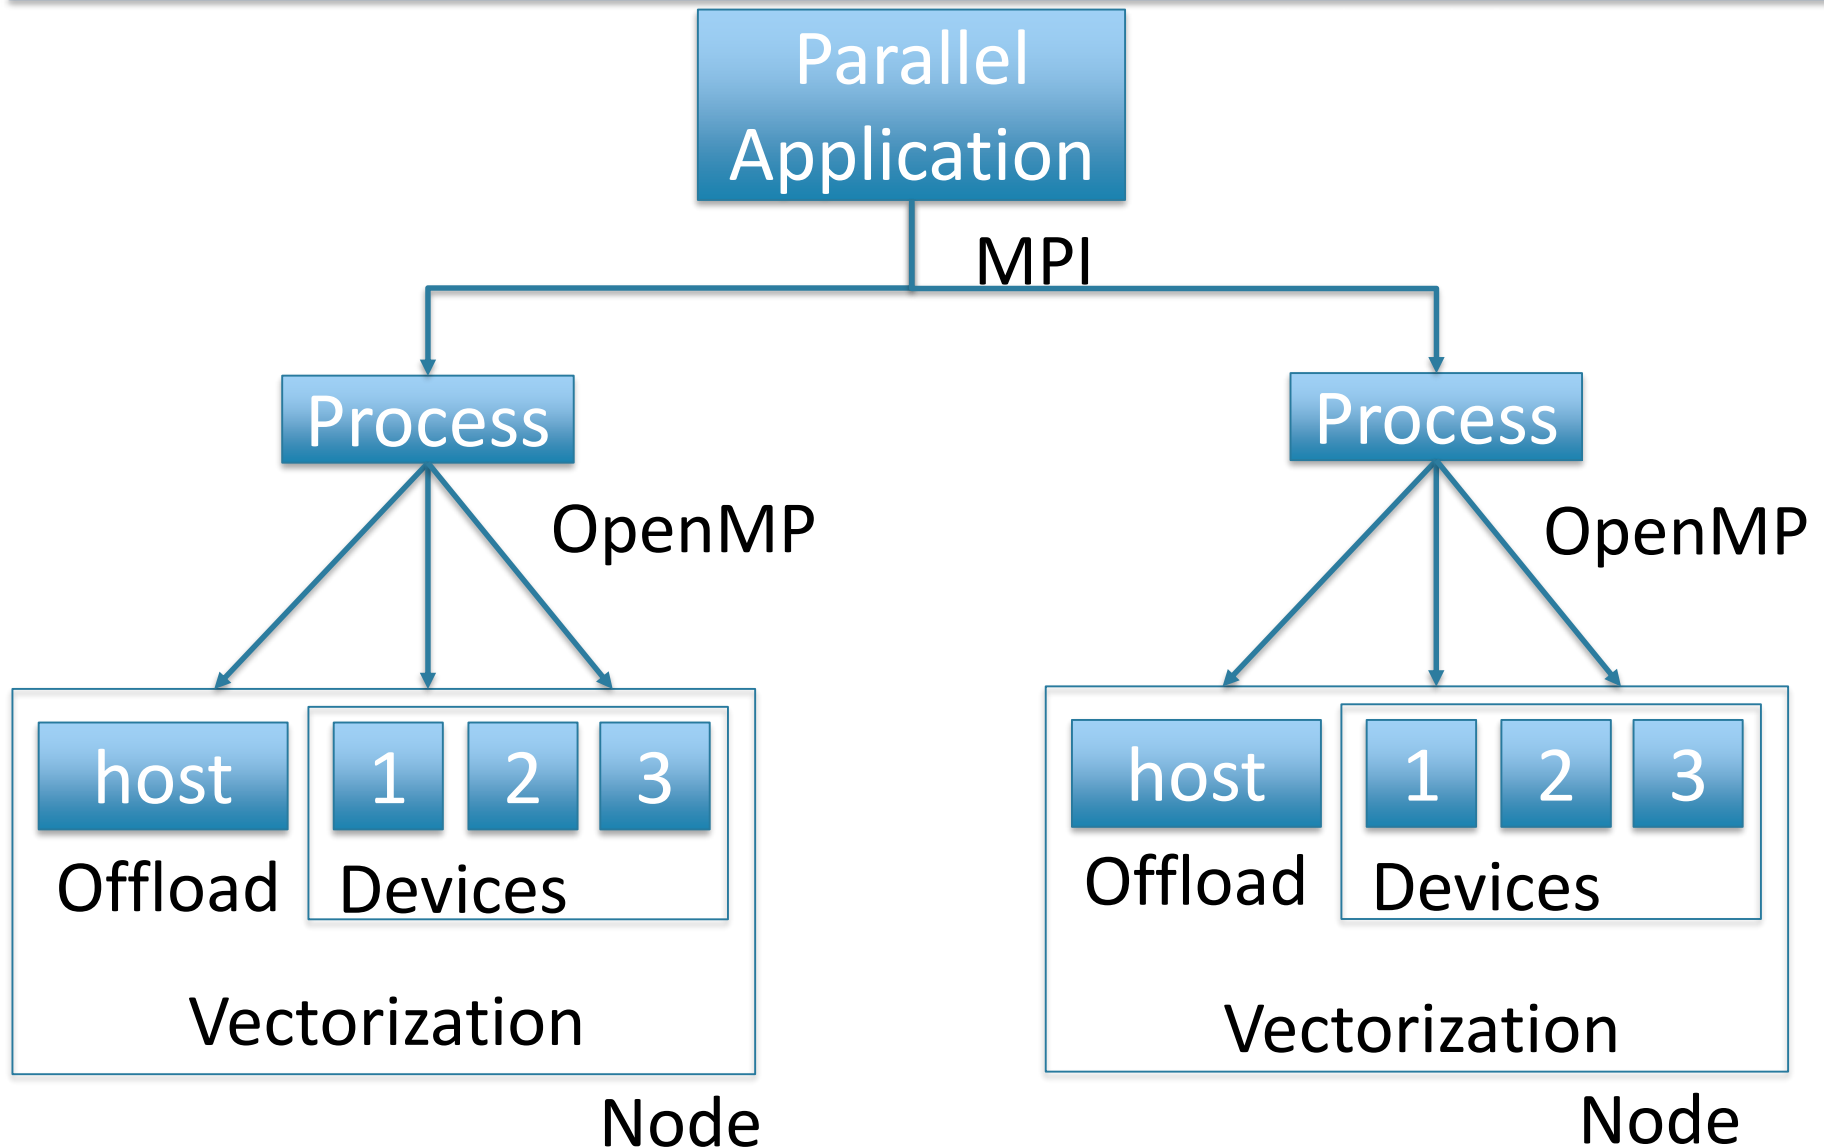
```
mpirun -n 8 ./myapp
```



**Infiniband:** computer-networking communications standard used in high-performance computing

- Very high throughput;
- Very low latency.

# Combining Levels






# Agenda

- Parallel Architectures
- Parallelism Levels
- **Xeon and Xeon Phi**
- Heterogeneous Cluster @ NCC
- Offload Programming
- Optimization Examples
- Hands-on


# Intel Xeon and Intel® Xeon Phi™ Overview

## Intel® Multicore Architecture




- ❖ Foundation of HPC Performance
- ❖ Suited for full scope of workloads
- ❖ Focus on fast single core/thread performance with “moderate” number of cores

## Intel® Many Integrated Core Architecture




- ❖ Performance and performance/watt optimized for highly parallelized compute workloads
- ❖ IA extension to Manycore
- ❖ Many cores/threads with wide SIMD

# Intel Xeon Architecture Overview



- Socket: mechanical component that provides mechanical and electrical connections between a microprocessor and a printed circuit board (PCB).
- QPI (Intel QuickPath Interconnect): high speed, packetized, point-to-point interconnection, that stitch together processors in distributed shared memory and integrated I/O platform architecture.




# Knights Landing (KNL)

Over 3 TF DP peak


Full Xeon ISA compatibility through AVX-512  
~3x single-thread vs. compared to Knights  
Corner

Up to 16GB high-bandwidth on-package  
memory (MCDRAM)  
Exposed as NUMA node  
~500 GB/s sustained BW

Up to 72  
cores  
2D mesh  
architecture



Tile



2x 512b VPU  
per core  
(Vector  
Processing  
Units)

6 channels  
DDR4  
Up to  
384GB

2 ports  
**Intel Omni-Path Fabric On-**  
**package**  
50 GB/s bi-directional

Based on Intel® Atom Silvermont processor with  
many HPC enhancements  
Deep out-of-order buffers  
Gather/scatter in hardware  
Improved branch prediction  
4 threads/core  
High cache bandwidth


# Cluster modes

## One single space address

### Hemisphere:

the tiles are divided into two parts called hemisphere


Node 0



### Quadrant:

tiles are divided into two parts called hemisphere or into four parts called quadrants

Node 0




# Cluster modes

Cache data are isolated in each sub numa domain


## SNC-2:

the tiles are  
divided into two  
Numa Nodes




## SNC-4:

the tiles are  
divided into two  
Numa Nodes



## Integrated On-Package Memory Usage Models




Split Options:  
25/75% or 50/50%

Cache Model	Flat Model	Hybrid Model
Hardware automatically manages the MCDRAM as a “L3 cache” between CPU and ext DDR memory	Manually manage how the app uses the integrated on-package memory and external DDR for peak perf	Harness the benefits of both Cache and Flat models by segmenting the integrated on-package memory
<ul style="list-style-type: none"><li>▪ App and/or data set is very large and will not fit into MCDRAM</li><li>▪ Unknown or unstructured memory access behavior</li></ul>	<ul style="list-style-type: none"><li>▪ App or portion of an app or data set that can be, or is needed to be “locked” into MCDRAM so it doesn’t get flushed out</li></ul>	<ul style="list-style-type: none"><li>▪ Need to “lock” in a relatively small portion of an app or data set via the Flat model</li><li>▪ Remaining MCDRAM can then be configured as Cache</li></ul>

# Vectorial processing Intel® Architecture

MMX	Intel® SSE	Intel® AVX / AVX2	Intel® MIC / AVX-512
Vector size: 64bit Data types: 8, 16 and 32 bit integers VL: 2,4,8	Vector size: 128bit Data types: 8,16,32,64 bit integers 32 and 64bit float VL: 2,4,8,16	Vector size: 256bit Data types: 32 and 64 bit floats VL: 4, 8, 16 Sample: Xi, Yi 32 bit int or float	Vector size: 512bit Data types: 32 and 64 bit integers 32 and 64bit float VL: 8,16



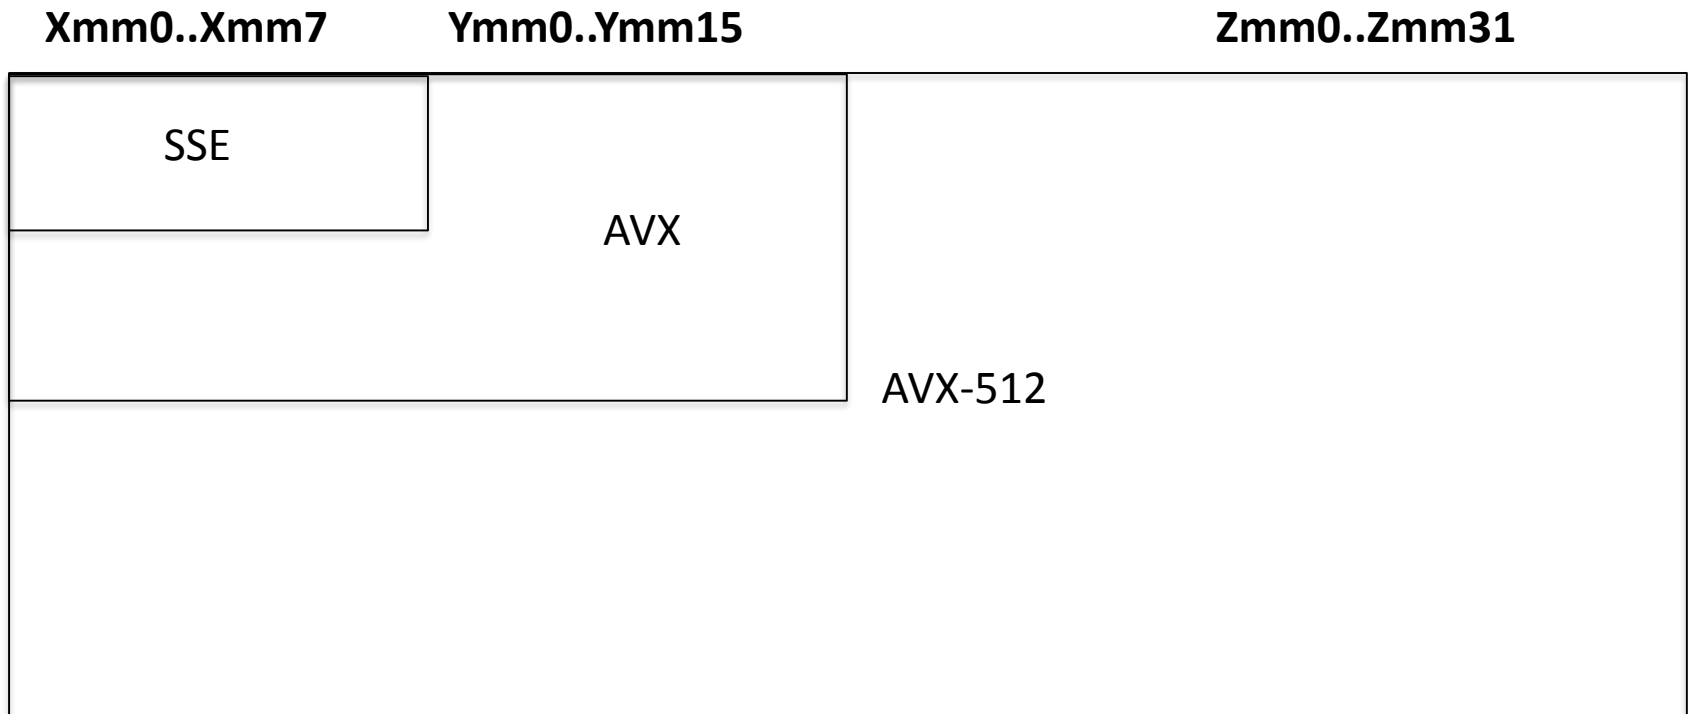
## AVX -512

Exponential & Reciprocal Instructions (ERI)  
 Prefetch Instructions (PFI)  
 Foundation instructions (FI)  
 Conflict Detection Instructions (CDI)  
 Byte & Word Instructions (BWI)  
 Double-/Quad-word Instructions (DQI)  
 Vector Length Extensions (VLE)



# AVX-512 - Foundation

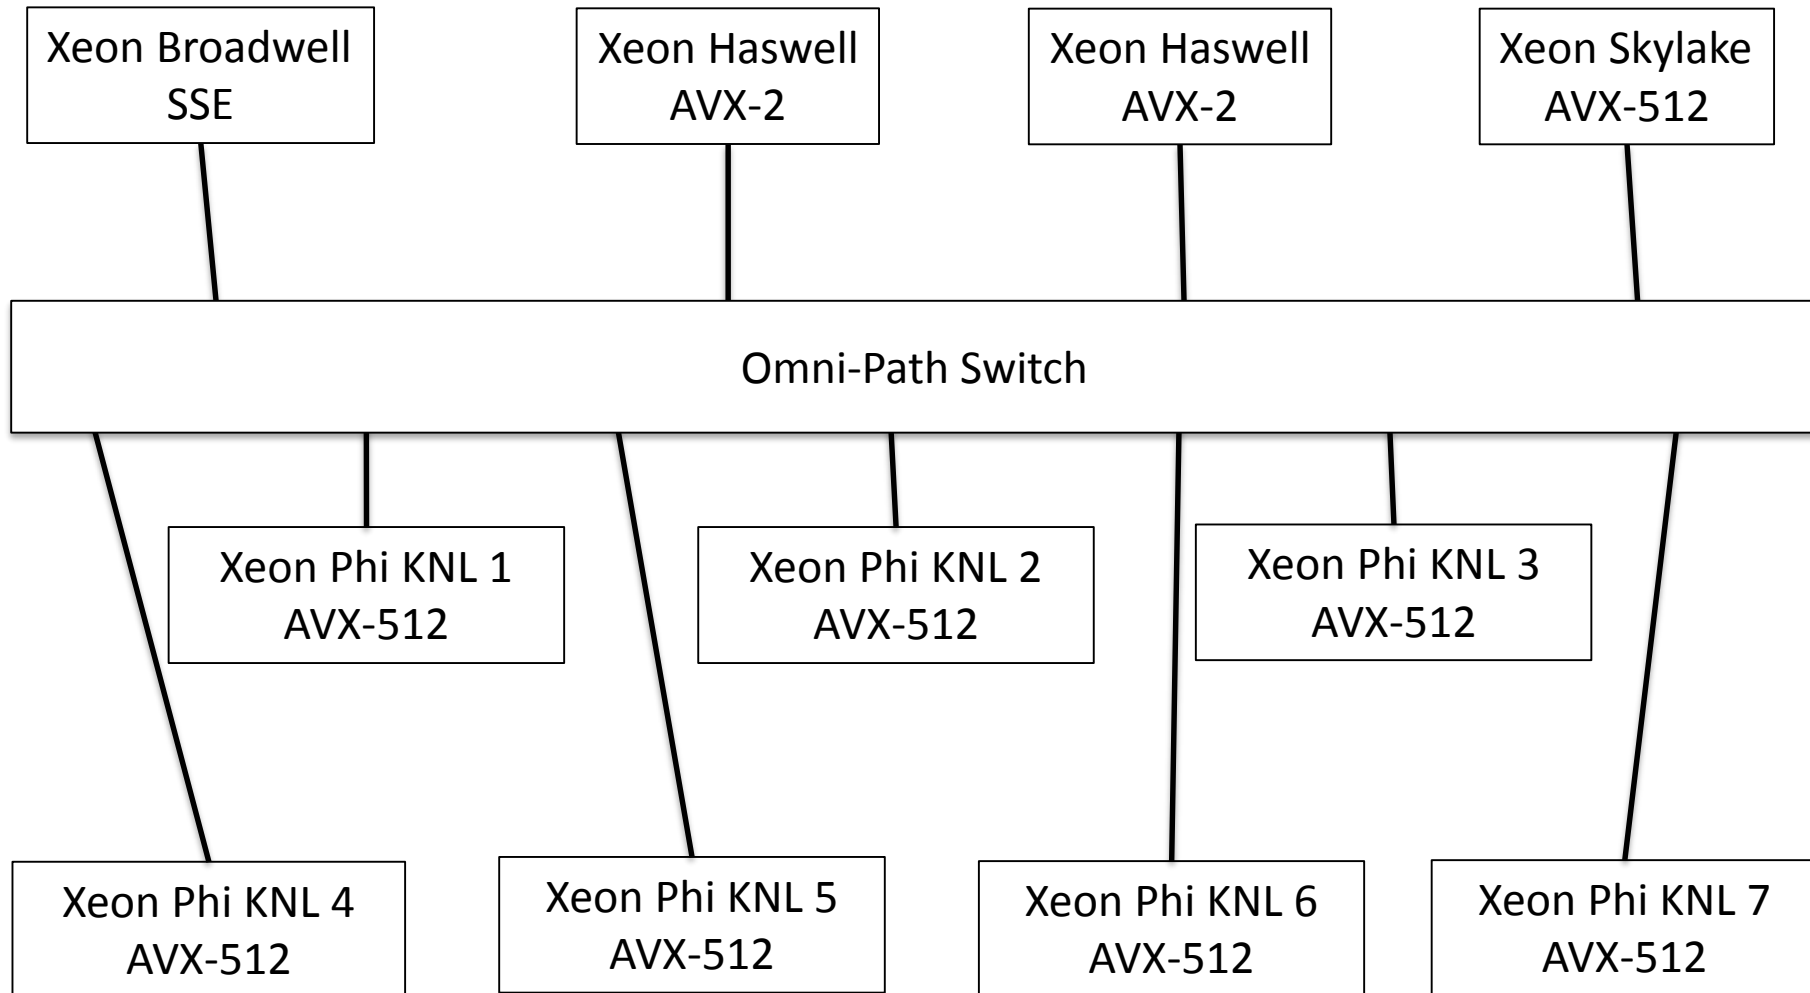
- Implements new version of AVX-2 functions using 512-bit registers (ZMM);



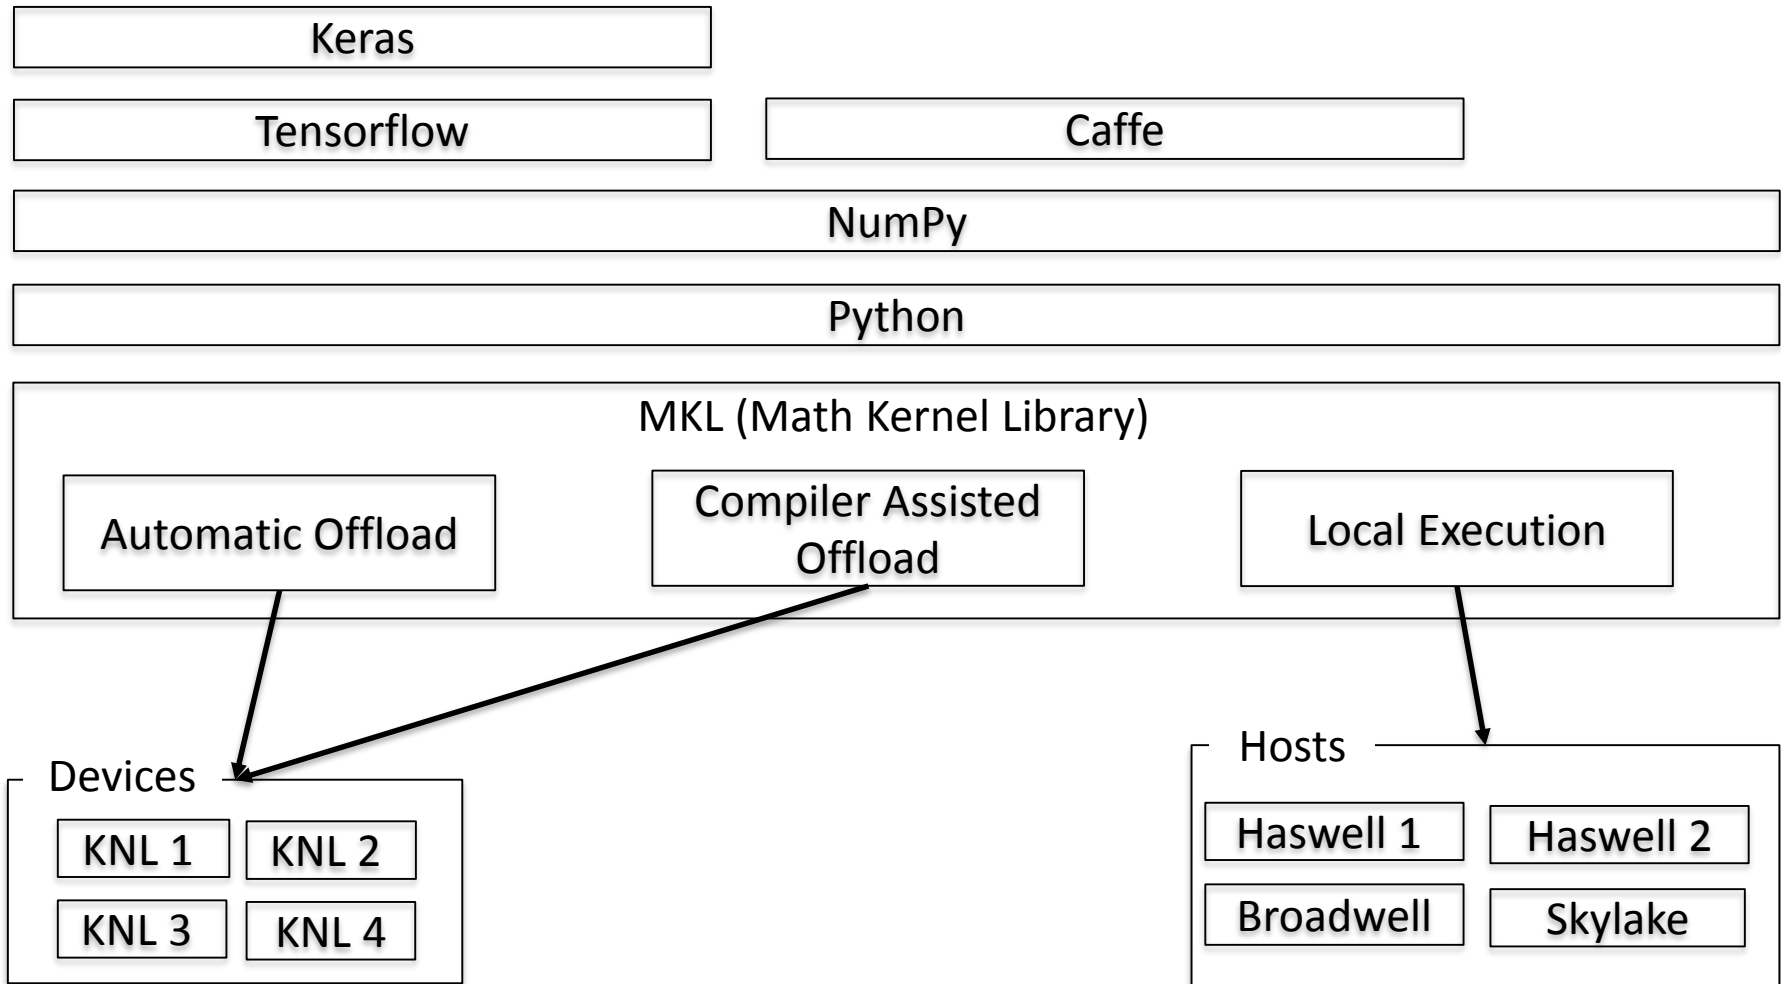
# Agenda

- Parallel Architectures
- Parallelism Levels
- Xeon and Xeon Phi
- **Heterogeneous Cluster @ NCC**
- Offload Programming
- Optimization Examples
- Hands-on

# Heterogeneous Cluster @ NCC



# Parallelism Model - MKL



# Parallelism Model – MPI x Offload

Hosts

Haswell 1

Haswell 2

Broadwell

Skylake

High Performance Network (Omni-Path)

Devices

KNL 1

KNL 2

KNL 3

KNL 4

Hosts

Haswell 1

Haswell 2

Broadwell

Skylake

High Performance Network (Omni-Path)

Devices

KNL 1

KNL 2

KNL 3

KNL 4

- **MPI (Message Passing Interface)**
  - High Performance Communication using MPI Functions
  - Every node send and receive messages

- **Offload**
  - One Host node can offload code to one or more devices
  - Hosts only send code regions for execution
  - Devices only execute code regions

# Offload Over Fabric

- Offload over Fabric (OOF) enables offloading from servers with Intel Xeon processor to servers with Intel Xeon Phi x200 processor within a high-speed network, such as Intel® Omni-Path Fabric.
- *OFFLOAD\_NODES* environment variable define nodes available for the offloading operation from the offload host.
  - *OFFLOAD\_NODES* is a comma-separated list of nodes' names or IP addresses.

# Offload Over Fabric

```
icc -qopenmp testHW.c -o testHW
```

```
export OFFLOAD_NODES=10.0.0.5  
./testHW
```

Running on host: phi01.ncc.unesp.br

Running on target: phi05

```
export OFFLOAD_NODES=10.0.0.6  
./testHW
```

Running on host: phi01.ncc.unesp.br

Running on target: phi06

# Agenda

- Parallel Architectures
- Parallelism Levels
- Xeon and Xeon Phi
- Heterogeneous Cluster @ NCC
- **Offload Programming**
- Optimization Examples
- Hands-on



# OpenMP 4.0 Offload

- **target:** transfers the control flow to the target device
  - Transfer is sequential and synchronous
  - Transfer clauses control data flow
- **target data:** creates a scoped device data environment
  - Does not include a transfer of control
  - Transfer clauses control data flow
  - The device data environment is valid through the lifetime of the target data region
- **target update:** request data transfers from within a target data region
- **omp declare target:** creates a structured-block of functions that can be offloaded.

# OpenMP 4.0 Offload Report

- **OFFLOAD REPORT:**
  - Measures the amount of time it takes to execute an offload region of code;
  - Measures the amount of data transferred during the execution of the offload region;
  - Turn on the report: `export OFFLOAD_REPORT=2`
- **[Var]** The name of a variable transferred and the direction(s) of transfer.
- **[CPU Time]** The total time measured for that offload directive on the host.
- **[MIC Time]** The total time measured for executing the offload on the target.
- **[CPU->MIC Data]** The number of bytes of data transferred from the host to the target.
- **[MIC->CPU Data]** The number of bytes of data transferred from the target to the host.

# Pragma omp declare target

- Creates a structured-block of functions that can be offloaded.
- Syntax
  - `#pragma omp declare target [clause[[,] clause],...]`  
*declaration of functions*
  - `#pragma omp end declare target`

# Pragma omp target

- Transfer control [and data] from host to device
- Syntax
  - `#pragma omp target [data] [clause[[,] clause],...]`  
*structured-block*
- Clauses
  - `device(scalar-integer-expression) :`
    - ❑ device to offload code;
  - `map(alloc | to | from | tofrom: list) :`
    - ❑ map variables to device;
  - `if(scalar-expr) :`
    - ❑ test an expression before offload:
      - o True executes on device;
      - o False executes on host;
  - `Nowait`
    - ❑ Execute the data transfer defined in map asynchronously;

# Pragma omp target

- Map clauses:
  - alloc : allocate memory on device;
  - to : transfer a variable from host to device;
  - from : transfer a variable from device to host;
  - tofrom :
    - ❑ transfer a variable from host to device before start execution;
    - ❑ transfer a variable from device to host after finish execution;

# Offloading - omp target

```
Int main() {  
  Printf("begin");  
  int N=25;  
  int b =2;  
  int l = 0;
```


*Offload:  
Copy variable:  
**N,b,l** and **a** to device*

```
#pragma omp target map(N,b,l,a)  
{  
  for (i=0; i<N; i++) a[i] = 2;  
  for (i=0; i<N; i++) a[i] = a[i] + b;  
}
```

```
for (i=0; i<N; i++)  
  printf("%d",a[i]);  
...  
return(0);  
}
```

Host

Device



# Pragma omp target example

```
#pragma omp target device(0) map(a[0:NUM][0:NUM])  
map(b[0:NUM][0:NUM]) map(c[0:NUM][0:NUM])  
{  
    #pragma omp parallel for collapse (2)  
    for(i=0; i<msize; i++) {  
        for(k=0; k<msize; k++) {  
            #pragma omp simd  
            for(j=0; j<msize; j++) {  
                c[i][j] = c[i][j] + a[i][k] * b[k][j];  
            }  
        }  
    }  
}
```

# Agenda

- Parallel Architectures
- Parallelism Levels
- Xeon and Xeon Phi
- Heterogeneous Cluster @ NCC
- Offload Programming
- **Optimization Examples**
- Hands-on



# N-Body Simulation

- An N-body simulation [1] aims to approximate the motion of particles that interact with each other according to some physical force;
- Used to study the movement of bodies such as satellites, planets, stars, galaxies, etc., which interact with each other according to the gravitational force;
- Newton's second law of motion can be used in a N-body simulation to define the bodies' movement.

[1] AARSETH, S. J. Gravitational n-body simulations. [S.l.]: Cambridge University Press, 2003. Cambridge Books Online.

# N-Body Algorithm

- Bodies struct:
  - 3 matrix represents velocity (x,y and z)
  - 3 matrix represents position (x,y and z)
  - 1 matrix represent mass
- A loop calculate temporal steps:
  - At each temporal step new velocity and position are calculated to all bodies according to a function that implements Newton's second law of motion

# N-Body - Parallel version (host only)

```
function Newton(step)
{
    #pragma omp for
    for each body[x] {
        #pragma omp simd
        for each body[y]
            calc force exerted from body[y] to body[x];
        calc new velocity of body[x]
    }
    #pragma omp simd
    for each body[x]
        calc new position of body[x]
}

Main() {
    for each temporal step
        Newton(step)
}
```

# N-Body - Parallel version (Load balancing)

- The temporal step loop remains sequential
- The N-bodies are divided among host and devices to be executed using Newton
- OpenMP offload pragmas are used to
  - Newton function offloading to devices
  - Transfer data (bodies) between host and devices

# N-Body - Parallel version (Load balancing)

```
function Newton(step, begin_body, end_body, deviceId)
{
    #pragma omp target device (deviceId) {
        #pragma omp for
        for each body[x] from subset(begin_body, end_body) {
            #pragma omp simd
            for each body[y] from subset(begin_body, end_body)
                calc force exerted from body[y] to body[x];
            calc new velocity of body[x]
        }
        #pragma omp simd
        for each body[x]
            calc new position of body[x]
    }
}
```

# N-Body - Parallel version (Load balancing)

for each temporal step


Divide the amount of bodies among host and devices;

```
#pragma omp parallel  
{  
    #pragma omp target data device ( tid ) to(bodies[begin_body:  
end_body])  
    {  
        Newton(step, begin_body, end_body, deviceId)  
        #pragma omp target update device ( tid ) (from:bodies)  
        #pragma omp barrier  
        #pragma omp target data device ( tid )  
to(bodies[begin_body: end_body])  
    }  
}
```

# N-Body

## N-Body Evaluation using offload over fabric (Xeon - Xeon Phi and Omni-Path)

Amount of Particles in Simulation: 30000



# Agenda

- Parallel Architectures
- Parallelism Levels
- Xeon and Xeon Phi
- Heterogeneous Cluster @ NCC
- Offload Programming
- Optimization Examples
- **Hands-on**



